

Threat Hunting and Remediation Using Agentic AI with Python and Microsoft Sentinel



This project demonstrates the design and implementation of an AI-assisted SOC workflow that supports threat hunting, investigation, and automated response using Microsoft Sentinel, Defender for Endpoint, and OpenAI. The agent translates natural language security concerns into KQL queries, analyzes enterprise telemetry, and executes controlled response actions with strict cost and safety guardrails. The result is a scalable, low-cost automation approach that enhances analyst efficiency without sacrificing oversight.

Introduction

This document walks through the process of building and testing an AI-powered agent designed to support threat hunting and remediation workflows within a Security Operations Center (SOC).

The goal of this project is to explore how AI agents can be integrated into a real-world SOC analyst workflow—specifically, how they can assist with querying security telemetry, interpreting results, and accelerating investigation efforts.

While technical steps are covered, the emphasis is placed on **how the solution supports real-world security workflows**, rather than on low-level implementation details.

Starting/Preparing the Environment

To begin, a local development environment must be set up to support Python-based development and secure communication with both OpenAI and Microsoft Azure services.

Step 1: Install Python and Set Up the IDE

The first requirement is installing **Python** on a Windows machine. Python serves as the foundation for the AI agent and is used to orchestrate API calls, process security logs, and format outputs.

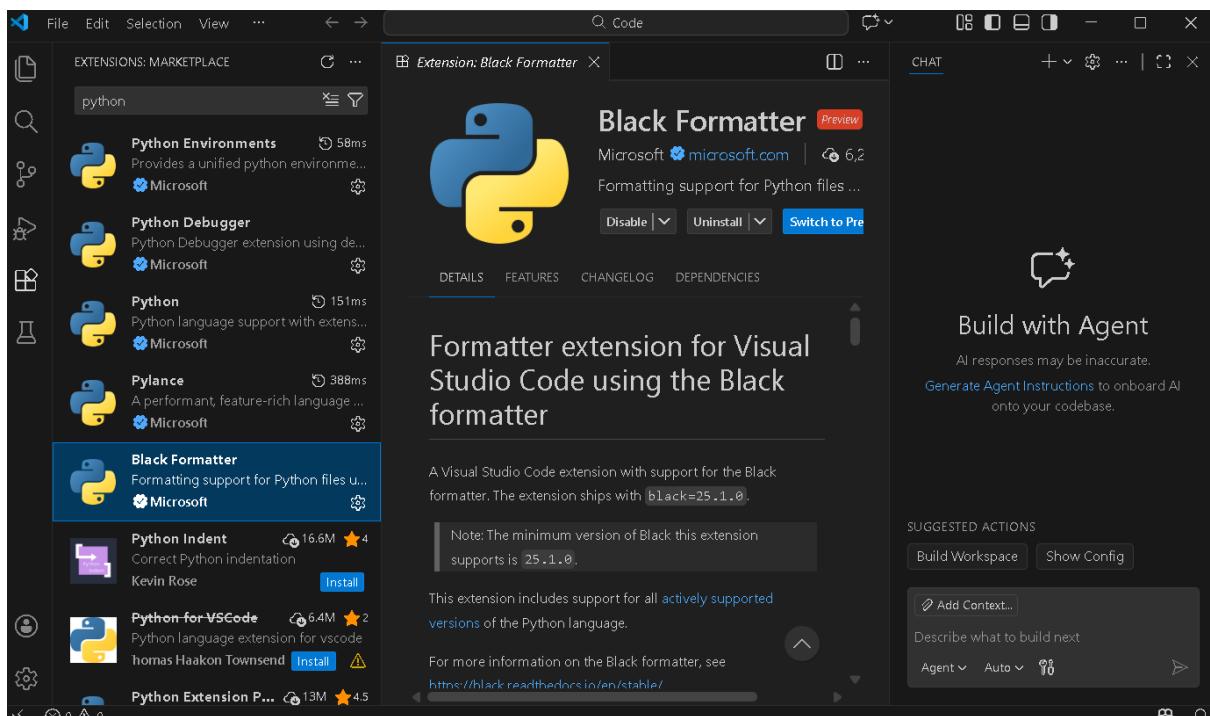
For development, **Visual Studio Code (VS Code)** was selected as the Integrated Development Environment (IDE) due to its lightweight nature and strong Python support.

To optimize VS Code for this project, the following extensions were installed from the VS Code Marketplace:

- **Python** – Enables Python syntax support, debugging, and execution

- **Black Formatter** – Automatically formats Python code to ensure consistency and readability

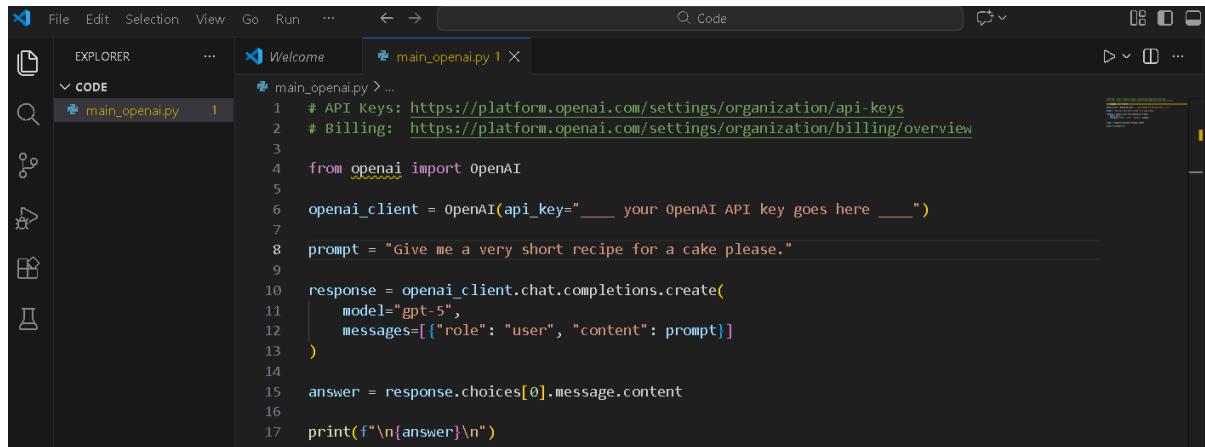
This setup provides a clean, efficient environment suitable for rapid experimentation and iteration.



Step 2: Preparing the AI Agent Template

With the environment configured, a basic Python template is created to serve as the starting point for the AI agent. This template establishes the structure needed to send prompts, receive responses, and later integrate external data sources.

At this stage, an **OpenAI API key** is generated, and the account is funded with a small amount of usage credit. This key allows the program to securely communicate with OpenAI's services. Usage costs are minimal and are discussed later in the project.



```
# API Keys: https://platform.openai.com/settings/organization/api-keys
# Billing: https://platform.openai.com/settings/organization/billing/overview

from openai import OpenAI

openai_client = OpenAI(api_key="_____ your OpenAI API key goes here _____")

prompt = "Give me a very short recipe for a cake please."

response = openai_client.chat.completions.create(
    model="gpt-3.5-turbo",
    messages=[{"role": "user", "content": prompt}]
)

answer = response.choices[0].message.content

print(f"\n{answer}\n")
```

Before inserting the API key into the code, it is necessary to install the OpenAI Python library. This is done using the following command:

```
PS C:\Users\laptop\Desktop\Code> py -m pip install openai
Collecting openai
  Downloading openai-2.16.0-py3-none-any.whl.metadata (29 kB)
Collecting anyio<5,>=3.5.0 (from openai)
  Downloading anyio-4.12.1-py3-none-any.whl.metadata (4.3 kB)
collecting distro<2,>=1.7.0 (from openai)
  Downloading distro-1.9.0-py3-none-any.whl.metadata (6.8 kB)
collecting httpx<1,>=0.23.0 (from openai)
  Downloading httpx-0.28.1-py3-none-any.whl.metadata (7.1 kB)
collecting jiter<1,>=0.10.0 (from openai)
  Downloading jiter-0.13.0-cp310-cp310-win_amd64.whl.metadata (5.3 kB)
collecting pydantic<3,>=1.9.0 (from openai)
  Downloading pydantic-2.12.5-py3-none-any.whl.metadata (90 kB)
```

Once installed, the API key is added to the appropriate location in the script.

Step 3: Verifying OpenAI API Connectivity

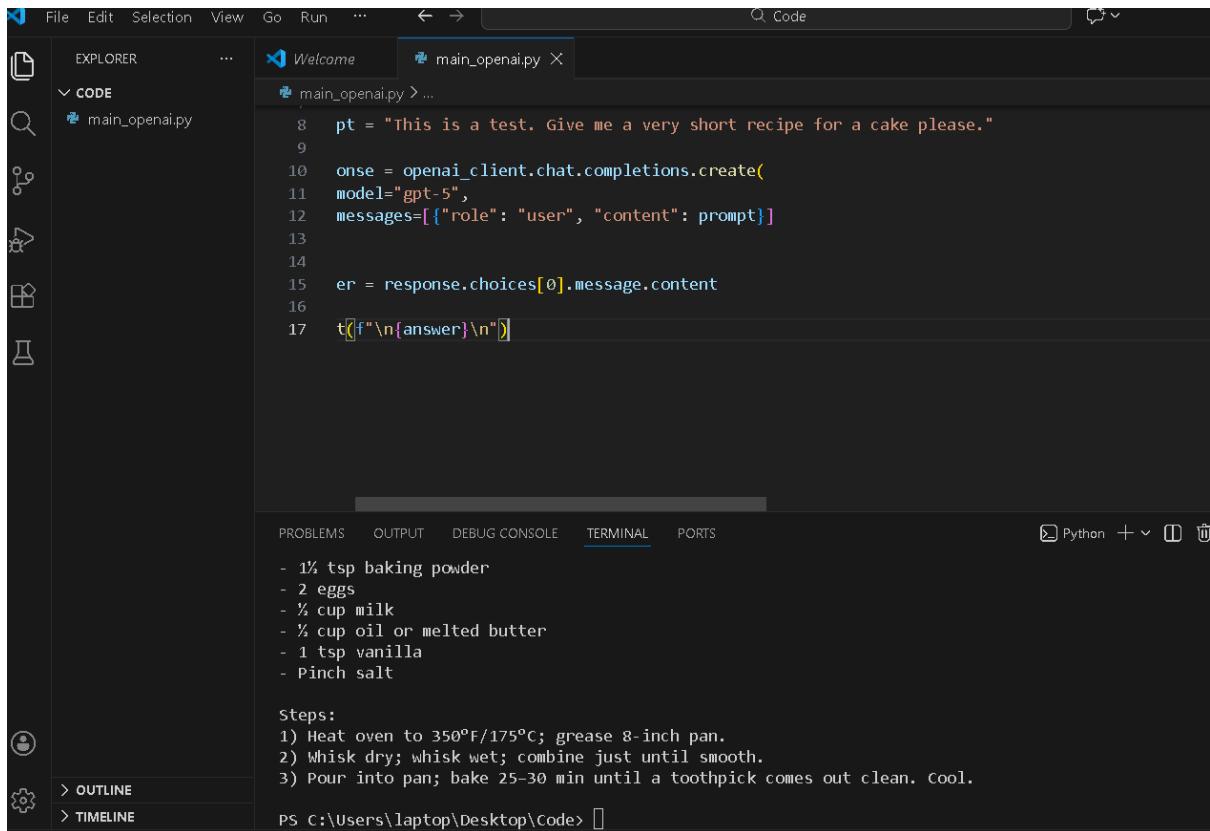
To confirm that the environment is configured correctly, a simple test prompt is sent to the OpenAI API:

“Give me a very short recipe for a cake.”

Receiving a valid response confirms that:

- The API key is working
- The program can successfully send requests
- Responses are being returned and processed correctly

This validation step ensures the AI agent's core communication mechanism is functioning before introducing additional complexity.



```

File Edit Selection View Go Run ... ← → ⌂ Code
EXPLORER ... Welcome main_openai.py ...
CODE main_openai.py > ...
8 pt = "This is a test. Give me a very short recipe for a cake please."
9
10 response = openai.Client.chat.Completions.create(
11     model="gpt-3.5-turbo",
12     messages=[{"role": "user", "content": prompt}]
13
14
15 er = response.choices[0].message.content
16
17 t(f"\n{answer}\n")

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

- 1½ tsp baking powder
- 2 eggs
- ½ cup milk
- ¼ cup oil or melted butter
- 1 tsp vanilla
- Pinch salt

Steps:
1) Heat oven to 350°F/175°C; grease 8-inch pan.
2) Whisk dry; whisk wet; combine just until smooth.
3) Pour into pan; bake 25–30 min until a toothpick comes out clean. Cool.

PS C:\Users\laptop\Desktop\Code>

Step 4: Introducing the Security Lab Environment

For realistic testing, this project uses a controlled lab environment known as **Cyber Range**. This environment simulates an enterprise organization that is continuously exposed to security incidents and breaches, providing realistic telemetry for threat hunting scenarios.

Using a lab environment allows the AI agent to operate against real-world-style data without risk to a production system.

The screenshot shows the Microsoft Azure Log Analytics workspace interface. On the left, there's a sidebar with various navigation options like Overview, Activity log, Access control (IAM), Tags, Diagnose and solve problems, and Logs (which is currently selected). The main area displays a query result table titled 'DeviceLogonEvents' with a limit of 'take 10'. The table has columns: TimeGenerated [UTC], AccountDomain, AccountName, and AccountSid. The results show log entries from April 2, 2026, at 06:47:34 UTC, for virtual users on sshd_6964 and sshd_5372. The table includes a 'Results' tab and a 'Chart' tab. At the bottom, there are buttons for 'Query details' and '1 - 7 of 10'.

Step 5: Connecting to Azure Log Analytics

The next step is integrating the AI agent with Microsoft Azure's **Log Analytics workspace**, which serves as the data backend for Microsoft Sentinel.

At this point, the project has:

- Verified connectivity to the OpenAI API
- Begun validating connectivity to Azure security data

The screenshot shows a code editor with a dark theme. The left sidebar shows files named 'log_analytics.py' and 'main_openai.py'. The main editor window contains Python code for connecting to Azure Log Analytics:

```

from datetime import timedelta
from azure.identity import DefaultAzureCredential
from azure.monitor.query import LogsQueryClient
import pandas as pd

LOG_ANALYTICS_WORKSPACE_ID = "i...z"

# Need Azure CLI: https://learn.microsoft.com/en-us/cli/azure/install-azure-cli-windows?view=...
log_analytics_client = LogsQueryClient(credential=DefaultAzureCredential())

hours_ago = 1

kql_query = f"""
DeviceLogonEvents
| take 10
...
"""

response = log_analytics_client.query_workspace(
    query=kql_query,
    timespan=timedelta(hours=hours_ago),
    workspace_id=LOG_ANALYTICS_WORKSPACE_ID
)
    
```

Below the code editor, there are tabs for 'PROBLEMS', 'OUTPUT', 'DEBUG CONSOLE', 'TERMINAL', and 'PORTS'. The status bar at the bottom shows 'Python' and other icons.

To enable communication with Azure services, several Python libraries are required:

`py -m pip install azure-identity`

`py -m pip install azure-monitor-query`

```
py -m pip install pandas
```

- **azure-identity** handles authentication
- **azure-monitor-query** enables Log Analytics queries
- **pandas** is used to structure and format data in a human-readable way

```
PS C:\Users\laptop\Desktop\Code> py -m pip install azure-identity
Collecting azure-identity
  Downloading azure_identity-1.25.1-py3-none-any.whl.metadata (88 kB)
Collecting azure-core>=1.31.0 (from azure-identity)
  Downloading azure_core-1.38.0-py3-none-any.whl.metadata (47 kB)
Collecting cryptography>=2.5 (from azure-identity)
  Downloading cryptography-46.0.4-cp311-abi3-win_amd64.whl.metadata (5.7 kB)
Collecting msal>=1.30.0 (from azure-identity)
  Downloading msal-1.34.0-py3-none-any.whl.metadata (11 kB)
Collecting msal-extensions>=1.2.0 (from azure-identity)
  Downloading msal_extensions-1.3.1-py3-none-any.whl.metadata (7.8 kB)
Requirement already satisfied: typing-extensions>=4.0.0 in C:\Users\laptop\AppData\Local\Python\pythoncore-3.14-64\Lib\site-packages (from azure-identity) (4.15.0)
```

```
PS C:\Users\laptop\Desktop\Code> py -m pip install azure-monitor-query
Collecting azure-monitor-query
  Downloading azure_monitor_query-2.0.0-py3-none-any.whl.metadata (20 kB)
Collecting isodate>=0.6.1 (from azure-monitor-query)
  Downloading isodate-0.7.2-py3-none-any.whl.metadata (11 kB)
Requirement already satisfied: azure-core>=1.30.0 in C:\Users\laptop\AppData\Local\Python\pythoncore-3.14-64\Lib\site-packages (from azure-monitor-query) (1.38.0)
Requirement already satisfied: typing-extensions>=4.6.0 in C:\Users\laptop\AppData\Local\Python\pythoncore-3.14-64\Lib\site-packages (from azure-monitor-query) (4.15.0)
Requirement already satisfied: requests>=2.21.0 in C:\Users\laptop\AppData\Local\Python\pythoncore-3.14-64\lib\site-packages (from azure-core>=1.30.0->azure-monitor-query) (2.32.5)
Requirement already satisfied: charset_normalizer<4,>=2 in C:\Users\laptop\AppData\Local\Python\pythoncore-3.14-64\lib\site-packages (from azure-monitor-query) (3.2.0)
```

Step 6: Authenticating with Azure

In addition to Python libraries, the **Azure CLI** must be installed on the local machine. The CLI is used to authenticate against an Azure tenant that has:

- Microsoft Sentinel enabled
- A user account with appropriate permissions to query logs

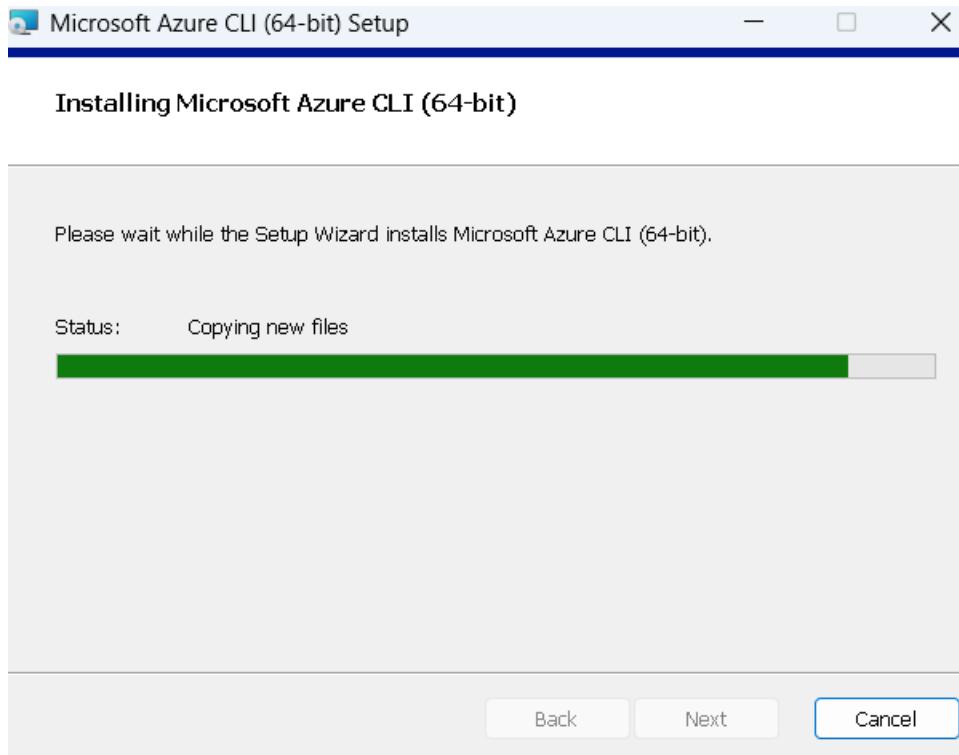
Latest version

Download and install the latest release of the Azure CLI. When the installer asks if it can make changes to your computer, select the "Yes" box.

[Latest MSI of the Azure CLI \(32-bit\)](#)

[Latest MSI of the Azure CLI \(64-bit\)](#)

If you have previously installed the Azure CLI, running either the 32-bit or 64-bit MSI will overwrite an existing installation.



Once installed, the Azure CLI is connected to the tenant, allowing the Python script to inherit authenticated access.

A screenshot of a code editor interface, likely Visual Studio Code, showing a Python file named "log_analytics.py". The code imports datetime, azure.identity, and azure.monitor.query, and defines a credential object. It then sets LOG_ANALYTICS_WORKSPACE_ID to a placeholder value and creates a LogAnalyticsClient object. The script queries the DeviceLogonEvents table over the last hour using Kusto Query Language (KQL).

```
File Edit Selection View Go Run ... ← → Q Code
EXPLORER ...
CODE
log_analytics.py >...
1 from datetime import timedelta
2 from azure.identity import DefaultAzureCredential
3 from azure.monitor.query import LogsQueryClient
4
5 LOG_ANALYTICS_WORKSPACE_ID = "c"
6
7 # Need Azure CLI: https://learn.microsoft.com/en-us/cli/azure/install-azure-cli-windows?view=az-devpreview
8 log_analytics_client = LogsQueryClient(credential=DefaultAzureCredential())
9
10 hours_ago = 1
11
12 kql_query = f"""
13 DeviceLogonEvents
14 | take 10
15 ...
16
17 response = log_analytics_client.query_workspace(
18     workspace_id=LOG_ANALYTICS_WORKSPACE_ID,
19     query=kql_query,
20     timespan=timedelta(hours=hours_ago)
21 )
```

Step 7: Testing Log Analytics Queries

With authentication complete, a test query is executed against the Log Analytics workspace. The script sends a **KQL (Kusto Query Language)** query to Azure, retrieves the results, and outputs them within the program.

At this stage, the raw output confirms that:

- The agent can successfully query Sentinel data
- Log data is being retrieved as expected

We can see the KQL query that the program will pass to the Azure Log analytics workspace instance.

```
kql_query = f'''  
DeviceLogonEvents  
| take 10  
'''
```

With the output being:

```
PS C:\Users\laptop\Desktop\Code> & C:/Users/laptop/AppData/Local/Python/pythoncore-3.14-64/python.exe c:/Users/laptop/Desktop/Code/log_analytics.py  
<azure.monitor.query._models.LogsTable object at 0x00000234B963AE40>  
[ 'TenantId', 'AccountDomain', 'AccountName', 'AccountSid', 'ActionType', 'AdditionalFields', 'AppGuardContainerId', 'DeviceId', 'DeviceName', 'FailureReason', 'InitiatingProcessAccountDomain', 'InitiatingProcessAccountName', 'InitiatingProcessAccountObjectId', 'InitiatingProcessAccoun tsid', 'InitiatingProcessAccountUpn', 'InitiatingProcessCommandLine', 'InitiatingProcessFileName', 'InitiatingProcessFolderPath', 'InitiatingProcessId', 'InitiatingProcessIntegrityLevel', 'InitiatingProcessMD5', 'InitiatingProcessParentFileName', 'InitiatingProcessParentId', 'InitiatingProcessSHA1', 'InitiatingProcessSHA256', 'InitiatingProcessTokenElevation', 'IsLocalAdmin', 'LogonId', 'LogonType', 'MachineGroup', 'Protocol', 'RemoteDeviceName', 'RemoteIP', 'RemoteIPTyp e', 'RemotePort', 'ReportId', 'Timestamp', 'TimeGenerated', 'InitiatingProcessParentCreationTim e', 'InitiatingProcessCreationTime', 'InitiatingProcessFileSize', 'InitiatingProcessVersionInfo CompanyName', 'InitiatingProcessVersionInfoFileDescription', 'InitiatingProcessVersionInfoInter nalFileName', 'InitiatingProcessVersionInfoOriginalFileName', 'InitiatingProcessVersionInfoProd
```

More complex queries are then tested to validate consistency and flexibility. These tests help confirm that the agent can handle increasingly realistic investigation scenarios.

```
kql_query = f'''  
DeviceLogonEvents  
| take 10  
| order by TimeGenerated desc  
| project TimeGenerated, AccountName, ActionType, DeviceName, RemoteIP  
'''
```

```
PS C:\Users\laptop\Desktop\Code> & C:/Users/laptop/AppData/Local/Python/pythoncore-3.14-64/python.exe c:/Users/laptop/Desktop/Cod e/log_analytics.py  
<azure.monitor.query._models.LogsTable object at 0x0000021456F22900>  
[ 'TimeGenerated', 'AccountName', 'ActionType', 'DeviceName', 'RemoteIP']  
[ 663221, tzinfo=<isodate.tzinfo.Utc object at 0x00000214568D2F90>], 'labuser', 'LogonSuccess', 'vm-final-lab-si', '10.0.0.8'], [dat etime.datetime(2026, 2, 4, 12, 16, 30, 492190, tzinfo=<isodate.tzinfo.Utc object at 0x00000214568D2F90>), 'labuser', 'LogonSuccess', 'vm-final-lab-si', '10.0.0.8'], [datetime.datetime(2026, 2, 4, 12, 16, 29, 694491, tzinfo=<isodate.tzinfo.Utc object at 0x0 0000214568D2F90>), 'labuser', 'LogonSuccess', 'vm-final-lab-si', '10.0.0.8'], [datetime.datetime(2026, 2, 4, 12, 16, 29, 584955, tzinfo=<isodate.tzinfo.Utc object at 0x00000214568D2F90>), 'labuser', 'LogonSuccess', 'vm-final-lab-si', '10.0.0.8'], [datetim e.datetime(2026, 2, 4, 12, 16, 29, 524344, tzinfo=<isodate.tzinfo.Utc object at 0x00000214568D2F90>), 'labuser', 'LogonSuccess', 'vm -final-lab-si', '10.0.0.8'], [datetime.datetime(2026, 2, 4, 12, 16, 29, 349639, tzinfo=<isodate.tzinfo.Utc object at 0x0000021456 8D2F90>), 'labuser', 'LogonSuccess', 'vm-final-lab-si', '10.0.0.8'], [datetime.datetime(2026, 2, 4, 12, 16, 29, 318328, tzinfo=<isod ate.tzinfo.Utc object at 0x00000214568D2F90>), 'labuser', 'LogonSuccess', 'vm-final-lab-si', '10.0.0.8'], [datetim e.datetime(2026, 2, 4, 12, 16, 29, 273904, tzinfo=<isodate.tzinfo.Utc object at 0x00000214568D2F90>), 'labuser', 'LogonSuccess', 'vm-final-la b-si', '10.0.0.8']]
```

Step 8: Improving Data Readability

Finally, additional code is introduced using **pandas** to convert the raw query results into structured tables. This step transforms machine-oriented output into a format that is easier for humans to interpret and analyze.

Improving readability is a critical step, as it mirrors how SOC analysts and hiring managers would expect insights to be presented—clear, structured, and actionable.

```
columns = table.columns
rows = table.rows
df = pd.DataFrame(rows,columns=columns)
df[ "TimeGenerated" ] = pd.to_datetime(df[ "TimeGenerated" ]).dt.strftime("%Y-%m-%d %H:%M:%S")
print(df)
```

```
PS C:\Users\laptop\Desktop\Code> & c:/Users/laptop/AppData/Local/Python/pythoncore-3.14-64/python.exe c:/Users/laptop/Desktop/Code/log_analytics.py
   TimeGenerated AccountName    ActionType          DeviceName      RemoteIP
0 2026-02-04 13:13:01      root  LogonFailed  linux-scan-erick.z  1.CX...  3.144
1 2026-02-04 13:12:11      dwm-1 LogonAttempted  vm-final-lab-pr      -
2 2026-02-04 13:12:09      umfd-0 LogonSuccess  vm-final-lab-pr      -
3 2026-02-04 13:12:09      umfd-0 LogonAttempted  vm-final-lab-pr      -
4 2026-02-04 13:12:09      umfd-1 LogonSuccess  vm-final-lab-pr  1.CX...  3.144
5 2026-02-04 13:12:03      root  LogonFailed  linux-scan-erick.z  1.CX...  3.144
6 2026-02-04 13:08:38      root  LogonFailed  linux-scan-erick.z  1.CX...  3.144
7 2026-02-04 13:07:47      root  LogonFailed  linux-scan-erick.z  1.CX...  3.144
8 2026-02-04 13:05:09      root  LogonFailed  linux-scan-erick.z  3.CX...  3.144
9 2026-02-04 12:59:53      root  LogonFailed  linux-scan-erick.z  8.CX...  8.144
PS C:\Users\laptop\Desktop\Code> []
```

We can be even more specific with our searches:

```
kql_query = f"""
BehaviorAnalytics
| project TimeGenerated, ActivityType, ActionType, UserName, EventSource, SourceIPAddress, SourceIPLocation
| where isnotempty (ActionType) and isnotempty (UserName)
| take 25
..."""
```

```
PS C:\Users\laptop\Desktop\Code> & c:/Users/laptop/AppData/Local/Python/pythoncore-3.14-64/python.exe c:/Users/laptop/Desktop/Code/log_analytics.py
   TimeGenerated ActivityType ...      SourceIPAddress      SourceIPLocation
0 2026-02-04 12:32:17      Logon ...           2                   76  medway, united states
1 2026-02-04 12:37:46      LogOn ...           2                   2  sydney, australia
2 2026-02-04 12:36:02 FailedLogon ...           2                   82  sydney, australia
3 2026-02-04 12:34:08      LogOn ...           2                   10  houston, united states
4 2026-02-04 12:32:18      LogOn ...           2                   76  medway, united states
5 2026-02-04 12:36:55      LogOn ...           2                   2  sydney, australia
6 2026-02-04 13:02:27      LogOn ...           2                   6  medway, united states
7 2026-02-04 13:01:21      LogOn ...           2                   76  medway, united states
8 2026-02-04 13:06:02      LogOn ...           2                   2  sydney, australia
9 2026-02-04 13:08:09      LogOn ...           2                   2  sydney, australia
10 2026-02-04 13:26:31     LogOn ...           2                   7  brighton, united kingdom
11 2026-02-04 13:23:33     LogOn ...           2                   4  london, united kingdom
In 36 Col 1  Spacing: 4
```

Adding the Agent

Step 1. The Groundwork for AI Functionality

At this stage, the focus shifts from environment setup to expanding the program's functionality and formally introducing the AI agent into the workflow.

Initially, the agent itself is not yet active within the program. Instead, the groundwork is laid by extending how data is retrieved from the Azure Log Analytics workspace. Additional tables are pulled into the application and exposed as variables within a dedicated function. This abstraction simplifies how the program requests data and maintains synchronization with the Log Analytics workspace.

```

log_analytics_functions.py > ...
1  from datetime import timedelta
2  from azure.identity import DefaultAzureCredential
3  from azure.monitor.query import LogsQueryClient
4  import pandas as pd
5
6  LOG_ANALYTICS_WORKSPACE_ID = "c"
7
8  TABLE_NAME = "AzureActivity" # DeviceLogonEvents, AzureNetworkAnalytics_CL, AzureActivity, SigninLogs
9
10 FIELDS = {
11     "DeviceLogonEvents": "TimeGenerated, AccountName, DeviceName, ActionType, RemoteIP, RemoteDeviceName",
12     "AzureNetworkAnalytics_CL": "TimeGenerated, FlowType_s, SrcPublicIPs_s, DestIP_s, DestPort_d, VM_s, AllowedInFlows_d",
13     "AzureActivity": "TimeGenerated, OperationNameValue, ActivityStatusValue, ResourceGroup, Caller, CallerIpAddress, Category",
14     "SigninLogs": "TimeGenerated, UserPrincipalName, OperationName, Category, ResultSignature, ResultDescription, AppDisp"
15 }
16
17 HOURS_AGO = 1
18
19 # Need Azure CLI: https://learn.microsoft.com/en-us/cli/azure/install-azure-cli-windows?view=azure-cli-latest&pivots=msi
20 log_analytics_client = LogsQueryClient(credential=DefaultAzureCredential())
21

```

```

PS C:\Users\laptop\Desktop\Code> & C:/Users/laptop/AppData/Local/Python/pythoncore-3.14-64/python.exe c:/Users/laptop/Desktop/Code/log_analytics_functions.py

AzureActivity
| project TimeGenerated, OperationNameValue, ActivityStatusValue, ResourceGroup, Caller, CallerIpAddress, Category
Timegenerated,OperationNameValue,ActivityStatusValue,ResourceGroup,Caller,CallerIpAddress,Category
2026-02-04 12:57:53.158505+00:00,MICROSOFT.COMPUTE/VIRTUALMACHINES/RUNCOMMAND/ACTION,Success,CYBER-RANGE-ADMIN-SOC,5
4
2026-02-04 13:01:14.917624+00:00,Microsoft.Advisor/recommendations/available/action,Active,$
    ,Microsoft.Advisor,0.0.0.0,
2026-02-04 12:55:37.734877+00:00,MICROSOFT.COMPUTE/VIRTUALMACHINES/WRITE,Start,$
    ,1
2026-02-04 12:55:37.984877+00:00,MICROSOFT.COMPUTE/VIRTUALMACHINES/WRITE,Accept,$
    ,4

```

By structuring data access in this way, specific fields can be predefined, filtered, and manipulated directly from the derived tables returned by the function. This design choice allows the program to remain **modular and adaptable**, making it easy to tailor the agent's inputs to different environments, datasets, or investigative use cases without rewriting core logic.

```

> def query_log_analytics(client, workspace_id, table, fields, timespan_hours_ago): ...
    fields_we_want_to_include = FIELDS[TABLE_NAME]

    results = query_log_analytics(
        client = log_analytics_client,
        workspace_id = LOG_ANALYTICS_WORKSPACE_ID,
        table = TABLE_NAME,
        fields = fields_we_want_to_include,
        timespan_hours_ago = HOURS_AGO
    )

# results = query_log_analytics(log_analytics_client, LOG_ANALYTICS_WORKSPACE_ID, TABLE_NAME, FIELDS[TABLE_NAME], HOURS_A

```

```

2026-02-04 14:03:34.888146+00:00,ExternalPublic,1      3|1|1|1|0|74|0,10.0.0.198,55615.0,cyber-range-admin-soc/linux-target-1,1.0,0,0,0,0,0,0,0
2026-02-04 14:03:34.888146+00:00,ExternalPublic,1      7|1|1|1|0|74|0,10.0.0.198,18145.0,cyber-range-admin-soc/linux-target-1,1.0,0,0,0,0,0,0,0
2026-02-04 14:03:34.888146+00:00,ExternalPublic,1      0|1|1|1|0|74|0,10.0.0.198,12988.0,cyber-range-admin-soc/linux-target-1,1.0,0,0,0,0,0,0,0
2026-02-04 14:03:34.888146+00:00,ExternalPublic,1      3|1|1|1|0|60|0,10.0.0.198,4009.0,cyber-range-admin-soc/linux-target-1,1.0,0,0,0,0,0,0,0
2026-02-04 14:03:34.888146+00:00,ExternalPublic,1      178|1|1|1|0|74|0 :   89|1|1|1|0|74|0,10.0.0.198,1961.0,cyber-range-admin-soc/linux-t
arget-1,2.0,0,0,0,0,0,0,0
2026-02-04 14:03:34.888146+00:00,ExternalPublic,1      73|1|1|1|0|60|0,10.0.0.198,9491.0,cyber-range-admin-soc/linux-target-1,1.0,0,0,0,0,0,0,0
2026-02-04 14:03:34.888146+00:00,ExternalPublic,206    3|1|1|1|0|74|0,10.0.0.198,101.0,cyber-range-admin-soc/linux-target-1,1.0,0,0,0,0,0,0,0
2026-02-04 14:03:34.888146+00:00,ExternalPublic,1      1|0|1|1|0|106|0,10.0.0.198,17184.0,cyber-range-admin-soc/linux-target-1,0.0,0,0,0,0,0,0,0
2026-02-04 14:03:34.888146+00:00,ExternalPublic,1      8|1|1|1|0|74|0,10.0.0.198,1234.0,cyber-range-admin-soc/linux-target-1,1.0,0,0,0,0,0,0,0
2026-02-04 14:03:34.888146+00:00,ExternalPublic,1      1|1|1|1|0|74|0,10.0.0.198,1521.0,cyber-range-admin-soc/linux-target-1,1.0,0,0,0,0,0,0,0
fin.
PS C:\Users\laptop\Desktop\Code>

```

```
def query_log_analytics(client, workspace_id, table, fields, timespan_hours_ago):  
  
    kql_query = f'''  
    {table}  
    | where FlowType_s == "MaliciousFlow"  
    | project {fields}  
    ...
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```
2026-02-04 13:14:11.327089+00:00, MaliciousFlow,, 10.0.0.198, 3306, cyber-range-admin-soc/linux-target-1, 1, 0, 0, 0
2026-02-04 13:53:50.900632+00:00, MaliciousFlow,, 10.0.0.155, 8728, cyber-range-admin-soc/windows-target-1, 1, 0, 0, 0
2026-02-04 13:53:50.900632+00:00, MaliciousFlow,, 10.0.0.155, 3389, cyber-range-admin-soc/windows-target-1, 4, 0, 0, 0
2026-02-04 13:53:50.900632+00:00, MaliciousFlow,, 10.0.0.155, 804, cyber-range-admin-soc/windows-target-1, 1, 0, 0, 0
2026-02-04 13:53:50.900632+00:00, MaliciousFlow,, 10.0.0.155, 805, cyber-range-admin-soc/windows-target-1, 1, 0, 0, 0
2026-02-04 13:53:50.900632+00:00, MaliciousFlow,, 10.0.0.198, 22, cyber-range-admin-soc/linux-target-1, 1, 0, 0, 0
2026-02-04 13:53:50.900632+00:00, MaliciousFlow,, 10.0.0.198, 3128, cyber-range-admin-soc/linux-target-1, 1, 0, 0, 0
2026-02-04 13:53:50.900632+00:00, MaliciousFlow,, 10.0.0.198, 8728, cyber-range-admin-soc/linux-target-1, 1, 0, 0, 0
2026-02-04 13:53:50.900632+00:00, MaliciousFlow,, 10.0.0.198, 22, cyber-range-admin-soc/linux-target-1, 3, 0, 0, 0
2026-02-04 13:53:50.900632+00:00, MaliciousFlow,, 10.0.0.198, 803, cyber-range-admin-soc/linux-target-1, 1, 0, 0, 0
2026-02-04 13:53:50.900632+00:00, MaliciousFlow,, 10.0.0.198, 8728, cyber-range-admin-soc/linux-target-1, 1, 0, 0, 0
2026-02-04 13:53:50.900632+00:00, MaliciousFlow,, 10.0.0.198, 16379, cyber-range-admin-soc/linux-target-1, 1, 0, 0, 0
```

fin.
PS C:\Users\laptop\Desktop\Code> █

```
2026-02-04 14:03:00.020026+00:00,{ @l c.com,Sign-in activity,SignInLogs,SUCCESS,,Azur
re Portal,{ "1":{ "city":"Bolton", "state":"Bolton", "countryOrRegion":"GB", "geoCoordinates":{ "latitude":53.57849884033203, "longitude":-2.4298999309539795} }
2026-02-04 14:04:41.786253+00:00,{ @lc c.com,Sign-in activity,SignInLogs,SUCCESS,,Azur
re Portal,{ "3":{ "city":"Pembroke Pines", "state":"Florida", "countryOrRegion":"US", "geoCoordinates":{ "latitude":26.02906036376953, "longitude":-80.31037139892578 } }
2026-02-04 14:08:53.960522+00:00,{ @lc c.com,Sign-in activity,SignInLogs,SUCCESS,,Azur
re Portal,{ "2":{ "city":"Mcdonough", "state":"Georgia", "countryOrRegion":"US", "geoCoordinates":{ "latitude":33.44905090332031, "longitude":-84.16316986083984 } }
2026-02-04 13:46:40.338891+00:00,{ @lo c.com,Sign-in activity,SignInLogs,SUCCESS,,Azur
re Portal,{ "5":{ "city":"Wayne", "state":"New Jersey", "countryOrRegion":"US", "geoCoordinates":{ "latitude":40.9520110839844, "longitude":-74.25946807861328 } }}
```

fin.
PS C:\Users\Janton\Desktop\Codes> █

```
984        0      0      0      0      0      0      0      0      0      0  
2026-02-04 13:51:54.123108+00:00,FileCreated,mdc-test-cyn,WER.  
          .tmp.csv,svchost.exe -k WerSvcGroup  
2026-02-04 13:55:33.124994+00:00,FileCreated,linux-target-1,pg.  
pg.1.sh,/bin/sh /usr/bin/apt-key --quiet --readonly --keyring /usr/share/keyrings/ubuntu-archive-keyring.gpg verify --status-fd 3 /tmp/apt.sig.camVM /tmp  
/apt.data.us1XHH  
2026-02-04 13:55:33.284671+00:00,FileCreated,linux-target-1,pg.  
pg.1.sh,/bin/sh /usr/bin/apt-key --quiet --readonly --keyring /usr/share/keyrings/ubuntu-archive-keyring.gpg verify --status-fd 3 /tmp/apt.sig.IUnSCA /tmp  
/apt.data.chjEJ  
2026-02-04 13:55:32.897845+00:00,FileCreated,linux-target-1,pg.  
pg.1.sh,/bin/sh /usr/bin/apt-key --quiet --readonly --keyring /usr/share/keyrings/ubuntu-archive-keyring.gpg verify --status-fd 3 /tmp/apt.sig.G%7EL /tmp  
/apt.data.mLkHRD  
2026-02-04 13:55:32.630249+00:00,FileCreated,linux-target-1,pg.  
pg.1.sh,/bin/sh /usr/bin/apt-key --quiet --readonly --keyring /usr/share/keyrings/ubuntu-archive-keyring.gpg verify --status-fd 3 /tmp/apt.sig.8tTqCM /tmp  
/apt.data.vNldq  
2026-02-04 14:00:37.494791+00:00,FileCreated,vm-final-lab-pr,energy-report.html,C:\ProgramData\Microsoft\Windows\Power Efficiency Diagnostics\energy-repor  
t.html,taskhostw.exe GAEvents%{Arg0}
```

```
TABLE_NAME = "DeviceFileEvents" # DeviceLogonEvents, AzureNetworkAnalytics_CL, AzureActivity, SigninLogs, DeviceFileEvents

FIELDS = [
    "DeviceLogonEvents": "TimeGenerated, AccountName, DeviceName, ActionType, RemoteIP, RemoteDeviceName",
    "AzureNetworkAnalytics_CL": "TimeGenerated, FlowType_s, SrcPublicIPs_s, DestIP_s, DestPort_d, VM_s, AllowedInFlows_d, AllowedOutFlows_d",
    "AzureActivity": "TimeGenerated, OperationNameValue, ActivityStatusValue, ResourceGroup, Caller, CallerIpAddress, Category",
    "SigninLogs": "TimeGenerated, UserPrincipalName, OperationName, Category, ResultSignature, ResultDescription, AppDisplayName, IPAddress, Location",
    "DeviceFileEvents": "TimeGenerated, ActionType, DeviceName, FileName, FilePath, InitiatingProcessCommandLine"
]
```

Step 2. Introducing AI Functionality to the program

With the data layer in place, AI capabilities are then integrated into the program. Initial testing focuses on **cost awareness and model selection**, ensuring the solution remains practical for ongoing use.

The program is designed to support multiple AI models, allowing the user to specify which model should be used at runtime. Cost-efficient models are defined explicitly within the code, giving flexibility to balance performance and operational expense.

For this implementation, **GPT-5** is selected as the active model.

```
tokens.py > ...
1  # API Keys: https://platform.openai.com/settings/organization/api-keys
2  # Billing: https://platform.openai.com/settings/organization/billing/overview
3  from openai import OpenAI
4  import logshort as logs
5  import tiktoken
6  import json
7  from colorama import init, Fore, Style
8
9  # Stuff you need: logshort.py, logslong.py (optional), pip install tiktoken, colorama
10
11 init(autoreset=True)
12
13 MAX_OUTPUT_TOKENS = 1_000
14 DEFAULT_MODEL = "gpt-5" * gpt-5, gpt-4.1, gpt-5-mini, gpt-4.1-nano
15 current_model = DEFAULT_MODEL
16
17 # Models: https://platform.openai.com/docs/models/compare
18 models = {
19     "gpt-5": {"max_window_tokens": 400_000, "max_output_tokens": 128_000, "cost_per_million_input": 1.25, "cost_per_million_output": 1.25},
20     "gpt-4.1": {"max_window_tokens": 1_047_576, "max_output_tokens": 32_768, "cost_per_million_input": 1.00, "cost_per_million_output": 1.00},
21     "gpt-5-mini": {"max_window_tokens": 400_000, "max_output_tokens": 128_000, "cost_per_million_input": 0.25, "cost_per_million_output": 0.25},
22     "gpt-4.1-nano": {"max_window_tokens": 1_000_000, "max_output_tokens": 32_768, "cost_per_million_input": 0.10, "cost_per_million_output": 0.10}
23 }
24
25 def calculate_chat_cost(prompt_tokens, cost_per_million) -> float:
26     cost = (prompt_tokens / 1_000_000) * cost_per_million
27     return round(cost, 6)
28
```

This modular approach to model selection allows the program to evolve over time, making it straightforward to switch models or adjust configurations as requirements change.

Testing the Agent

Step 1. Evaluating Inputs and Managing Cost

With the AI model integrated, the next step is providing the agent with data to analyze.

```
1 security_logs = r''
2 TimeGenerated [UTC] AccountDomain AccountName ActionType DeviceName FileName InitiatingProcessCommandLine ProcessCommandLine
3 8/14/2025, 8:58:40.903 PM edr-andres irlab14# ProcessCreated edr-andres userinit.exe winlogon.exe userinit.exe
4 8/14/2025, 8:58:41.081 PM edr-andres irlab14# ProcessCreated edr-andres explorer.exe userinit.exe Explorer.EXE
5 8/14/2025, 8:58:43.651 PM edr-andres irlab14# ProcessCreated edr-andres svchost.exe services.exe svchost.exe -k ClipboardSvcGroup -p DllHost.exe /Process
6 8/14/2025, 8:58:44.444 PM edr-andres irlab14# ProcessCreated edr-andres svchost.exe svchost.exe -k DcomLaunch -p DllHost.exe /Process
7 8/14/2025, 8:58:45.041 PM nt authority system ProcessCreated edr-andres MicrosoftEdgeUpdate.exe services.exe "MicrosoftEdgeUpdate.exe" /NetworkRestricted -p ctfmon.exe svchost.exe -k LocalSystemNetworkRestricted -p MicrosoftEdgeUpdate.exe "MicrosoftEdgeUpdate.exe" /c "Micro
8 8/14/2025, 8:58:46.853 PM edr-andres irlab14# ProcessCreated edr-andres MicrosoftEdgeUpdate.exe "MicrosoftEdgeUpdate.exe" /c "Micro
9 8/14/2025, 8:58:48.342 PM nt authority system ProcessCreated edr-andres MicrosoftEdgeUpdate.exe "MicrosoftEdgeUpdate.exe" /c "Micro
10 8/14/2025, 8:58:55.547 PM nt authority system ProcessCreated edr-andres SearchFilterHost.exe SearchIndexer.exe /Embedding "Search
11 8/14/2025, 8:58:56.783 PM edr-andres irlab14# ProcessCreated edr-andres StartMenuExperienceHost.exe svchost.exe -k DcomLaunch -p "S
12 8/14/2025, 8:58:57.190 PM nt authority system ProcessCreated edr-andres wermgr.exe "MicrosoftEdgeUpdate.exe" /svc "wermgr.exe" "-outi
13 8/14/2025, 8:58:57.843 PM nt authority system ProcessCreated edr-andres svchost.exe services.exe svchost.exe -k WerSvcGroup
14 8/14/2025, 8:58:58.371 PM edr-andres irlab14# ProcessCreated edr-andres RuntimeBroker.exe svchost.exe -k DcomLaunch -p RuntimeBro
15 8/14/2025, 8:59:00.412 PM edr-andres irlab14# ProcessCreated edr-andres SearchApp.exe svchost.exe -k DcomLaunch -p "SearchApp.exe"
16 8/14/2025, 8:59:01.692 PM edr-andres irlab14# ProcessCreated edr-andres RuntimeBroker.exe svchost.exe -k DcomLaunch -p RuntimeBro
17 8/14/2025, 8:59:02.459 PM nt authority system ProcessCreated edr-andres MicrosoftEdgeUpdate.exe "MicrosoftEdgeUpdate.exe" /svc "Micro
18 8/14/2025, 8:59:04.854 PM edr-andres irlab14# ProcessCreated edr-andres SkypeApp.exe svchost.exe -k DcomLaunch -p "SkypeApp.exe"
19 8/14/2025, 8:59:05.254 PM edr-andres irlab14# ProcessCreated edr-andres backgroundTaskHost.exe svchost.exe -k DcomLaunch -p "backg
20 8/14/2025, 8:59:05.256 PM edr-andres irlab14# ProcessCreated edr-andres skypeBackgroundHost.exe svchost.exe -k DcomLaunch -p "Skype
21 8/14/2025, 8:59:09.331 PM edr-andres irlab14# ProcessCreated edr-andres backgroundTaskHost.exe svchost.exe -k DcomLaunch -p "Backg
22 8/14/2025, 8:59:09.575 PM edr-andres irlab14# ProcessCreated edr-andres dllhost.exe svchost.exe -k DcomLaunch -p DllHost.exe /Proce
23 8/14/2025, 8:59:09.651 PM edr-andres irlab14# ProcessCreated edr-andres taskhostw.exe svchost.exe -k netsvcs -p taskhostw.exe Sync
24 8/14/2025, 8:59:10.547 PM edr-andres irlab14# ProcessCreated edr-andres mbsync.exe svchost.exe -k DcomLaunch -p mbsync.exe -Embedd
25 8/14/2025, 8:59:10.922 PM edr-andres irlab14# ProcessCreated edr-andres wlrmrd.exe winlogon.exe -c -s 0 -f 0 -t Empty -m Empty -a i
26 8/14/2025, 8:59:12.193 PM nt authority system ProcessCreated edr-andres cmd.exe "CollectGuestLogs.exe" -Mode:ga -FileName:D:\CollectGu
27 8/14/2025, 8:59:13.101 PM edr-andres irlab14# ProcessCreated edr-andres RuntimeBroker.exe svchost.exe -k DcomLaunch -p RuntimeBro
28 8/14/2025, 8:59:14.292 PM edr-andres irlab14# ProcessCreated edr-andres SearchProtocolHost.exe searchIndexer.exe /Embedding "Search
29 8/14/2025, 8:59:15.518 PM edr-andres irlab14# ProcessCreated edr-andres RuntimeBroker.exe svchost.exe -k DcomLaunch -p RuntimeBro
```

```
tokens.py X main_openai.py logsshortpy
tokens.py > ...
48
49 prompt = f"You are a threat hunter. Any anomalies or potential breaches?\n{logs.security_logs}"
50
51 prompt_messages = [{"role": "user", "content": prompt}]
52
53 estimated_prompt_token_count = count_message_tokens(messages=prompt_messages, model=DEFAULT_MODEL)
54
55 # We can restrict max output tokens
56 max_tokens_used_in_response = MAX_OUTPUT_TOKENS
57 # max_tokens_used_in_response = models[current_model]["max_output_tokens"]
58 model_cost_per_million_tokens_input = models[current_model]["cost_per_million_input"]
59 model_cost_per_million_tokens_output = models[current_model]["cost_per_million_output"]
60
61 estimated_chat_cost_input = calculate_chat_cost(
62     prompt_tokens=estimated_prompt_token_count,
63     cost_per_million=model_cost_per_million_tokens_input
64 )
65
66 estimated_chat_cost_output = calculate_chat_cost(
67     prompt_tokens=max_tokens_used_in_response,
68     cost_per_million=model_cost_per_million_tokens_output
69 )
70
71 estimated_total_chat_cost = estimated_chat_cost_input + estimated_chat_cost_output
72
73 choice = input(
74     f"Chat will cost approximately ${estimated_total_chat_cost:.4f}\n"
75     "Proceed? (y/n): "
76 ).strip().lower()
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
PS C:\Users\laptop\Desktop\Code>
```

Before any analysis is performed, the program evaluates the projected cost of the request. This ensures that each investigation is assessed for resource usage prior to execution.

```
PS C:\Users\laptop\Desktop\Code> & C:/Users/laptop/AppData/Local/Python/pythoncore-3.14-64/python.exe c:/Users/laptop/Desktop/Code/tokens.py
Chat will cost approximately $0.0821
Proceed? (y/n):
```

```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS Python +
PS C:\Users\laptop\Desktop\Code> & C:/Users/laptop/AppData/Local/Python/pythoncore-3.14-64/python.exe c:/Users/laptop/Desktop/Code/tokens.py
Chat will cost approximately $0.0421
Proceed? (y/n): y

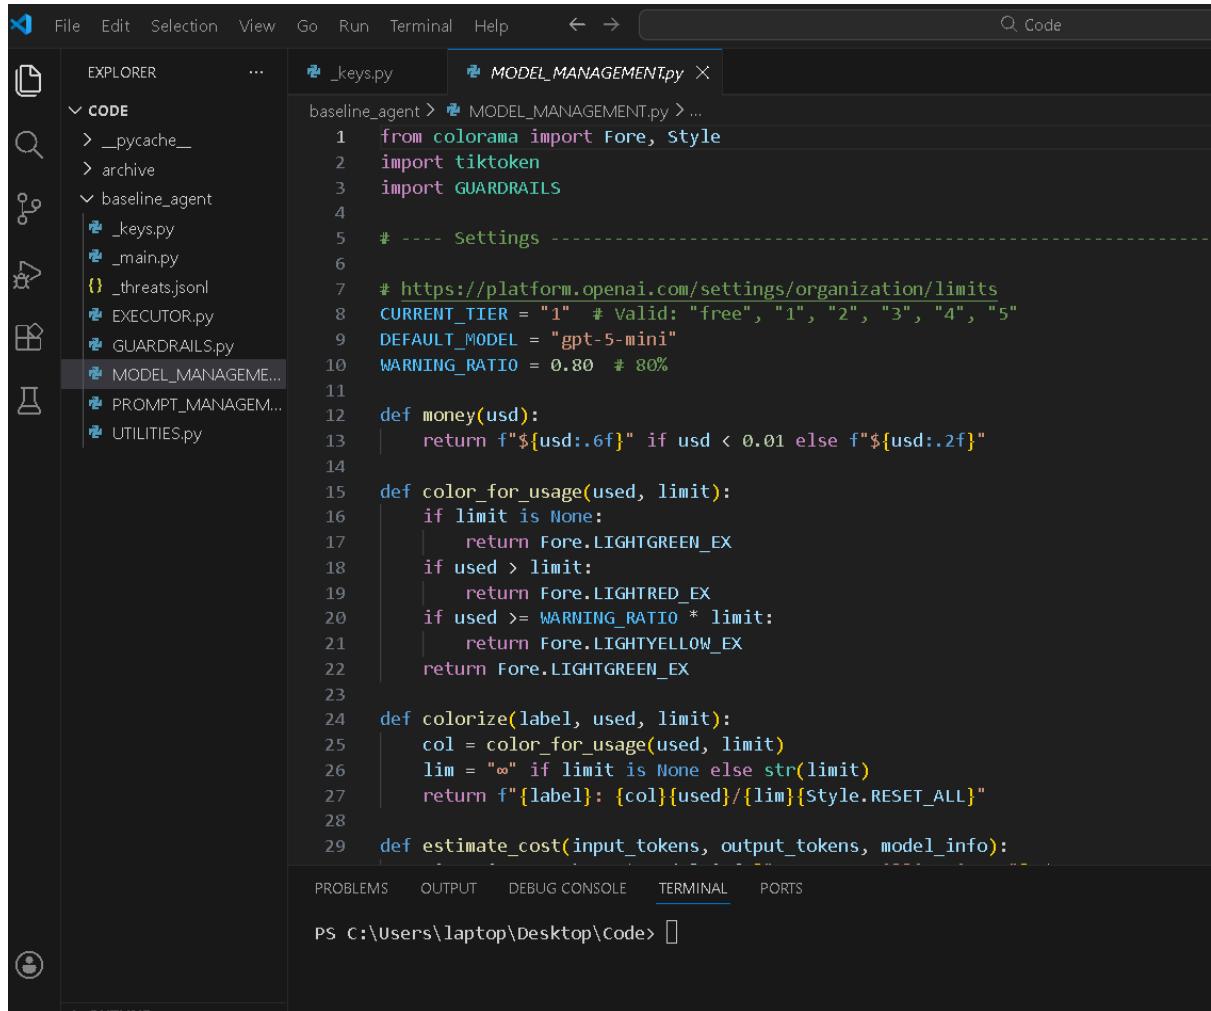
Estimated input tokens: 25693
Estimated output tokens: 1000
Estimated chat cost: $0.0421

Actual input tokens: 24547 [UNDER]
Actual output tokens: 1000 [UNDER]
Actual chat cost: $0.0407 [UNDER]

fin.
PS C:\Users\laptop\Desktop\Code>
```

When token usage is maximized, the cost of a single investigation is approximately **\$0.08**. To prevent unexpected or unnecessary spending, guardrails are implemented by explicitly defining a maximum token limit. In this case, a limit of **1,000 tokens** is later applied—an

amount that proves sufficient for the intended use case while keeping costs consistently low.



The screenshot shows a code editor interface with the following details:

- File Menu:** File, Edit, Selection, View, Go, Run, Terminal, Help.
- Search Bar:** Code
- Explorer:** Shows a file tree with the following structure:
 - EXPLORER
 - CODE
 - > __pycache__
 - > archive
 - baseline_agent
 - _keys.py
 - _main.py
 - _threats.jsonl
 - EXECUTOR.py
 - GUARDRAILS.py
 - MODEL_MANAGEMENT.py
 - PROMPT_MANAGEMENT.py
 - UTILITIES.py
- Code Editor:** Displays the content of `MODEL_MANAGEMENT.py`. The code includes imports for `colorama`, `tiktoken`, and `GUARDRAILS`. It defines variables for `CURRENT_TIER` (set to "1"), `DEFAULT_MODEL` (set to "gpt-3.5-mini"), and `WARNING_RATIO` (set to 0.80). It contains functions for `money(usd)`, `color_for_usage(used, limit)`, `colorize(label, used, limit)`, and `estimate_cost(input_tokens, output_tokens, model_info)`.
- Terminal:** Shows the command `PS C:\Users\laptop\Desktop\Code>`

With cost controls in place, the AI agent functionality is fully enabled within the program.

Step 2. Agent Behaviour and Threat Identification

The AI is configured to operate under a defined **SOC Analyst role**, established through system-level instructions within the program. This ensures the agent's behaviour aligns with how a human analyst would approach an investigation—prioritizing context, relevance, and risk.

The agent connects to the Azure Log Analytics workspace through the Azure CLI, while OpenAI handles the natural language processing and reasoning behind each request. This separation allows the program to securely access enterprise telemetry while leveraging AI for interpretation and decision support.

To validate functionality, the following test prompt is issued:

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

Agentic SOC Analyst at your service! What would you like to do?

I'm worried someone has breached our tenant. Can you search for suspicious sign ins in the last 24 hours?[]

```

"I'm worried someone has breached our tenant. Can you search for suspicious sign-ins in the last 24 hours?"

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

Model limits and estimated total cost:

gpt-4.1-nano | input limit: 169606/1047576 | rate_limit: 169606/200000 | out_max: 32768 | est_cost: $0.02
gpt-4.1      | input limit: 169606/1047576 | rate_limit: 169606/30000 | out_max: 32768 | est_cost: $0.17
gpt-5-mini   | input limit: 169606/272000 | rate_limit: 169606/200000 | out_max: 128000 | est_cost: $0.04 <-- (current)
gpt-5      | input limit: 169606/272000 | rate_limit: 169606/30000 | out_max: 128000 | est_cost: $0.22

 Safe: input limit: 169606/272000 is within the input limit for gpt-5-mini.
⚠️ WARNING: rate_limit: 169606/200000 is at least 80% of the TPM rate limit for gpt-5-mini.

Continue with 'gpt-5-mini'? (Enter to continue / type a model name / 'list'):
Selected model is valid: gpt-5-mini

Initiating cognitive threat hunt against targete logs...

Cognitive hunt complete. Took 134.82 seconds and found 3 potential threat(s)!

Press [Enter] or [Return] to see results. []

```

Ln 8, Col 35 S

```

I'm worried someone has breached our tenant. Can you search for suspicious sign ins in the last 24 hours?

Deciding log search parameters based on user request...

Log search parameters finalized:
Table Name: SigninLogs
Time Range: 24 hour(s)
Fields:    TimeGenerated, UserPrincipalName, OperationName, Category, ResultSignature, ResultDescription, AppDisplayName, IPAddress, LocationDetails

Rationale for log search parameters selection:
User requested suspicious sign-ins across the tenant in the last 24 hours. SigninLogs is the high-signal source for AAD sign-ins. 24h window covers recent activity. Fields chosen capture user, time, operation, result, app, IP, and location for detection of unusual sign-ins (impossible travel, risky sign-ins, failed MFA). No specific UPN or device provided so query is tenant-wide.

Validating Tables and Fields...
Fields and tables have been validated and comply with the allowed guidelines.

Constructed KQL Query:
SigninLogs
| where UserPrincipalName startswith ""
| project TimeGenerated, UserPrincipalName, OperationName, Category, ResultSignature, ResultDescription, AppDisplayName, IPAddress, LocationDetails

Querying Log Analytics Workspace ID: '6...' c'...
1105 record(s) returned.

Building threat hunt prompt/instructions...

Model limits and estimated total cost:

gpt-4.1-nano | input limit: 169606/1047576 | rate_limit: 169606/200000 | out_max: 32768 | est_cost: $0.02
gpt-4.1      | input limit: 169606/1047576 | rate_limit: 169606/30000 | out_max: 32768 | est_cost: $0.17

```

Ln 8, Col 35 Spaces:4 UTF-8 LF {} Python 88 3

From this prompt, the program automatically derives intent and generates an appropriate **KQL query** based on the requested investigation whilst also managing cost. Multiple query fields—such as time range, authentication context, and sign-in characteristics—are dynamically created by the AI.

```

Building threat hunt prompt/instructions...

Model limits and estimated total cost:

gpt-4.1-nano | input limit: 169606/1047576 | rate_limit: 169606/200000 | out_max: 32768 | est_cost: $0.02
gpt-4.1 | input limit: 169606/1047576 | rate_limit: 169606/30000 | out_max: 32768 | est_cost: $0.17
gpt-5-mini | input limit: 169606/272000 | rate_limit: 169606/200000 | out_max: 128000 | est_cost: $0.04 <-- (current)
gpt-5 | input limit: 169606/272000 | rate_limit: 169606/30000 | out_max: 128000 | est_cost: $0.22

☒ Safe: input limit: 169606/272000 is within the input limit for gpt-5-mini.
⚠WARNING: rate_limit: 169606/200000 is at least 80% of the TPM rate limit for gpt-5-mini.

Continue with 'gpt-5-mini'? (Enter to continue / type a model name / 'list'):
Selected model is valid: gpt-5-mini

Initiating cognitive threat hunt against target logs...

Cognitive hunt complete. Took 134.82 seconds and found 3 potential threat(s)!

Press [Enter] or [Return] to see results.

===== Potential Threat #1 =====

Title: Sign-in attempts blocked from IP flagged as malicious ( 1 )

Description: Azure AD blocked sign-in attempts because the source IP was known malicious. Multiple blocked failures targeting Azure Active Directory PowerShell occurred from : 1 within a short time window – consistent with automated credential abuse or targeted probing.

Confidence Level: High

MITRE ATT&CK Info:
Tactic: Credential Access

===== Potential Threat #1 =====

Title: Sign-in attempts blocked from IP flagged as malicious ( 1 )

Description: Azure AD blocked sign-in attempts because the source IP was known malicious. Multiple blocked failures targeting Azure Active Directory PowerShell occurred from : 1 within a short time window – consistent with automated credential abuse or targeted probing.

Confidence Level: High

MITRE ATT&CK Info:
Tactic: Credential Access
Technique: T1110
Sub-technique: T1110.001 (Password Spraying / Brute Force)
ID: T1110, T1110.001
Description: Attacker attempted sign-ins (likely brute-force / scripted) from an IP with known malicious activity; Azure blocked the attempts.

Log Lines:
- 2026-02-03 22:01:11.612459+00:00,0          @0lc      "Sign-in activity,SignInLogs,FAILURE, 
Sign-in was blocked because it came from an IP address with malicious activity,Azure Active Directory PowerShell,1        1,"{"city":"Tokyo",...}"
- 2026-02-03 22:07:44.804976+00:00,6          @0lc      "Sign-in activity,SignInLogs,FAILURE, 
Sign-in was blocked because it came from an IP address with malicious activity,Azure Active Directory PowerShell,1        1,"{"city":\\"Tokyo",...}"

Indicators of Compromise:
- IP: 1
- Target app: Azure Active Directory PowerShell
- Target UPN: c@l0c.com
- Timestamps: 2026-02-03 22:01:11Z, 2026-02-03 22:07:44Z

Tags:
- credential access
- password spraying

Ln 8, Col 35 Spaces:4 UTF-8 LF {} Python ⚙ 3

===== Potential Threat #2 =====

Title: Multiple invalid-credential failures from diverse IPs – potential password-spray activity

Description: Several different accounts show 'Error validating credentials due to invalid username or password' failures across multiple remote IPs and regions within a short timeframe. Pattern of many failed authentications across accounts/IPs is consistent with password-spray or opportunistic brute forcing.

Confidence Level: Medium

MITRE ATT&CK Info:
Tactic: Credential Access
Technique: T1110
Sub-technique: T1110.001 (Password Spraying)
ID: T1110, T1110.001
Description: Repeated failed authentication events across accounts/IPs indicating attempted credential guessing/password-spray.

Log Lines:
- 2026-02-04 11:22:12.297038+00:00,          @0lc      "Sign-in activity,SignInLogs,FAILURE, 
Error validating credentials due to invalid username or password.,Azure Active Directory PowerShell,3        3,"{"city":"Los Angeles",...}"
- 2026-02-04 12:32:34.046899+00:00,6          @0lc      "Sign-in activity,SignInLogs,FAILURE, 
Error validating credentials due to invalid username or password.,ACOM Azure Website,1        1,"{"city\":\"Lisboa",...}"
- 2026-02-04 01:01:31.908902+00:00,3          @0lc      "Sign-in activity,SignInLogs,FAILURE, 
Error validating credentials due to invalid username or password.,Azure Portal,9        9,"{"city\":\"Providence",...}"

Indicators of Compromise:
- IPs: 37, 16, 19, 5, 24
- Sample UPNs: 0, f7..., b6, 1..., 3, ...
- Messages: 'Error validating credentials due to invalid username or password.'

```

The agent then analyzes the data returned from the Log Analytics query, identifies potential threats, and records the findings in a structured **JSON file** for later review.

```
constructed KQL Query:  
signinLogs  
| where UserPrincipalName startswith ""  
| project TimeGenerated, UserPrincipalName, OperationName, Category, ResultSignature, ResultDescription, AppDisplayName, IPAddress, LocationDetails
```

Logged 3 threats to _threats.jsonl.

```
_keyspy MODEL_MANAGEMENT.py _threatsjsonl ×  
_threats.jsonl  
1 {"title": "Sign-in attempts blocked from IP flagged as malicious (1)", "description": "Azure AD blocked sign-in attempts be  
2 {"title": "Multiple invalid-credential failures from diverse IPs – potential password-spray activity", "description": "Several different  
3 {"title": "Repeated Azure CLI sign-ins from a single foreign IP with MFA interrupts (1)", "description": "A single identity  
4 |
```

Each identified event is enriched with:

- A mapped attack category or technique
- Relevant investigative or response guidelines
- A confidence level to assist analyst judgment

This approach mirrors how SOC analysts' triage and assess alerts, enabling faster review while maintaining transparency and analyst oversight.

Step 3: Automated Threat Response

With the core investigation workflow in place, the final phase of the project introduces **automated threat response** capabilities through **Microsoft Defender for Endpoint (MDE)**. This enhancement transforms the agent from a passive analysis tool into an active participant in incident response.

To support this expanded capability, additional guardrails are implemented to ensure actions are deliberate, controlled, and auditable. These safeguards are designed to prevent overreaction, limit scope, and ensure that automated responses occur only when specific criteria are met as seen by the results.

Agentic SOC Analyst at your service! What would you like to do?

I'm worried that windows-target-1 may have been maliciously logged into in the last few days.□

Step 4: Automated Investigation and Response Workflow

At this stage, the agent can take investigative findings and translate them directly into response actions.

When the AI identifies suspicious or malicious activity, it now:

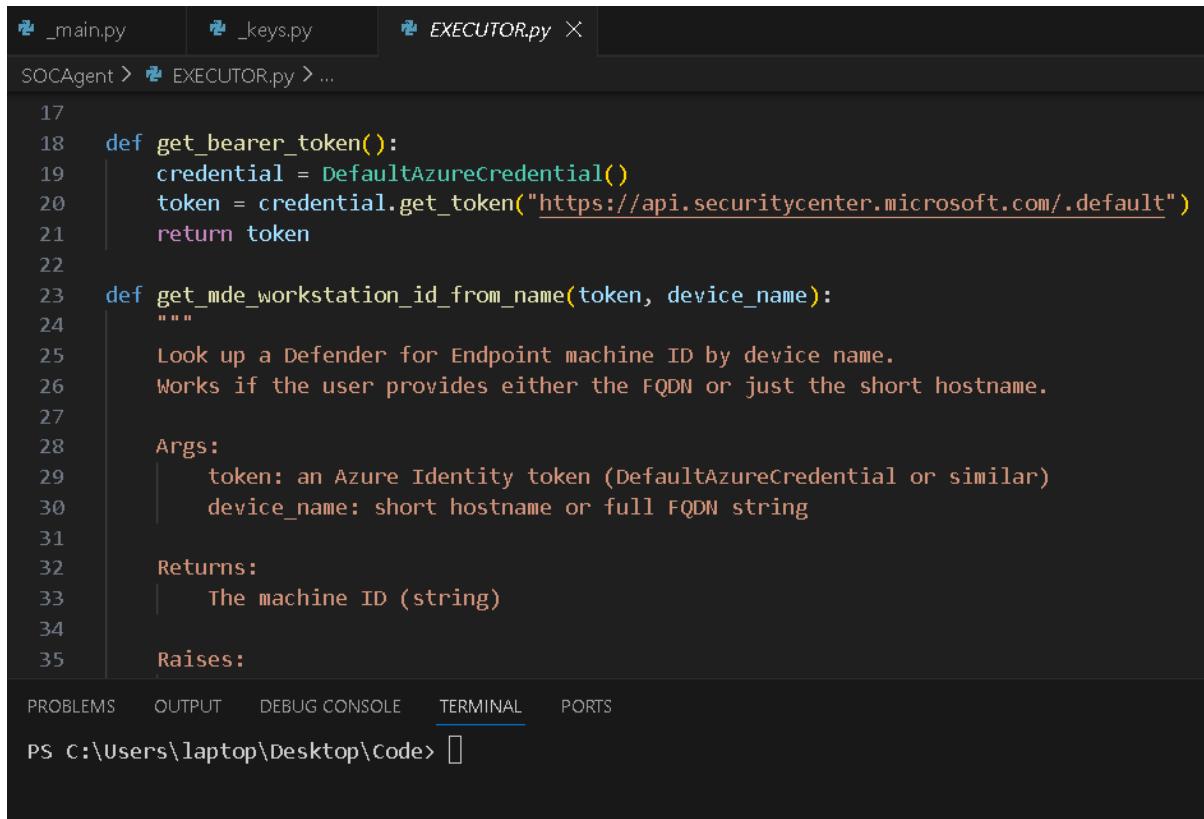
- Generates a **custom KQL query** based on the investigation context
- Searches historical and real-time records within the Log Analytics workspace
- Correlates affected users, sessions, and endpoints

- Determines whether a response action is warranted

```
Query context and metadata:
Table Name: DeviceLogonEvents
Time Range: 96 hour(s)
Fields: TimeGenerated, AccountName, DeviceName, ActionType, RemoteIP, RemoteDeviceName
Device: windows-target-1
User Related: False
Host Related: True
NSG Related: False
Rationale:
User concerned about potential malicious logon to a specific Windows host. DeviceLogonEvents contains event-level logon activity (local/remote), so it's the most relevant table. chosen time range 96 hours (4 days) per instruction when no timeframe provided. Fields selected to show timestamp (TimeGenerated), account used (AccountName), affected host (DeviceName), type of logon action (ActionType), remote IP (RemoteIP) and remote host name (RemoteDeviceName) to help identify suspicious remote access sources.
```

Constructed KQL Query:

```
DeviceLogonEvents
| where DeviceName startswith "windows-target-1"
| project TimeGenerated, AccountName, DeviceName, ActionType, RemoteIP, RemoteDeviceName
```



```

_main.py      _keys.py      EXECUTOR.py ×
SOCAgent > EXECUTOR.py > ...

17
18 def get_bearer_token():
19     credential = DefaultAzureCredential()
20     token = credential.get_token("https://api.securitycenter.microsoft.com/.default")
21     return token
22
23 def get_mde_workstation_id_from_name(token, device_name):
24     """
25         Look up a Defender for Endpoint machine ID by device name.
26         Works if the user provides either the FQDN or just the short hostname.
27
28     Args:
29         token: an Azure Identity token (DefaultAzureCredential or similar)
30         device_name: short hostname or full FQDN string
31
32     Returns:
33         The machine ID (string)
34
35     Raises:

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

PS C:\Users\laptop\Desktop\Code>

```

# If the machine is already isolated, don't isolate it again in the same session (wastes API calls)
if threat_confidence_is_high and (not machine_is_isolated):

    print(Fore.YELLOW + "[!] High confidence threat detected on host:" + Style.RESET_ALL, query_context["device_name"])
    print(Fore.LIGHTRED_EX + threat['title'])
    confirm = input(f"[{Fore.RED}{Style.BRIGHT}Would you like to isolate this VM? (yes/no): " + Style.RESET_ALL).strip().lower()

    if confirm.startswith("y"):
        machine_id = EXECUTOR.get_mde_workstation_id_from_name(
            token=token,
            device_name=query_context["device_name"]
        )
        machine_is_isolated = EXECUTOR.quarantine_virtual_machine(
            token=token,
            machine_id=machine_id
        )
        if machine_is_isolated:
            print(Fore.GREEN + "[+] VM successfully isolated." + Style.RESET_ALL)
            print(Fore.CYAN + "Reminder: Release the VM from isolation when appropriate at: " + Style.RESET_ALL + "https://security.microsoft.com")
    else:
        print(Fore.CYAN + "[i] Isolation skipped by user." + Style.RESET_ALL)

```

If malicious activity is confirmed and falls within predefined thresholds, the agent initiates an automated response through Microsoft Defender for Endpoint. In this implementation, the primary response action is **isolating affected virtual machines**, effectively containing the threat while further analysis can continue.

```

Logged 4 threats to _threats.jsonl.

[!] High confidence threat detected on host: windows-target-
High-volume, distributed failed logon attempts targeting privileged accounts on windows-target-
Would you like to isolate this VM? (yes/no): 

```

```

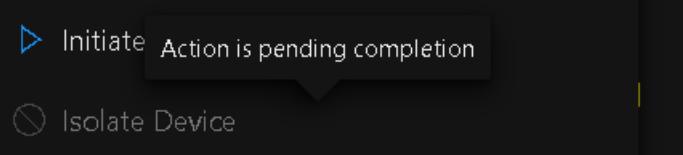
[!] High confidence threat detected on host: windows-target-
High-volume, distributed failed logon attempts targeting privileged accounts on windows-target-
Would you like to isolate this VM? (yes/no): yes
[+] VM successfully isolated.
Reminder: Release the VM from isolation when appropriate at: https://security.microsoft.com/
PS C:\Users\laptop\Desktop\Code> 

```

This mirrors real-world SOC escalation workflows, where containment is prioritized to limit impact while preserving evidence.

Step 5: Action Validation and Logging

Once a response action is issued, the agent monitors the status returned by MDE.



A “pending completion” status confirms that the action was successfully submitted and is in progress, providing immediate feedback that the containment step has been initiated.

Device isolation

Status
Device isolation pending
[Cancel action](#)

⌚ Device isolation submitted
Isolation via Python Agentic AI using DefaultAzureCredential
By [redacted]
on 5 Feb 2026 03:36:13

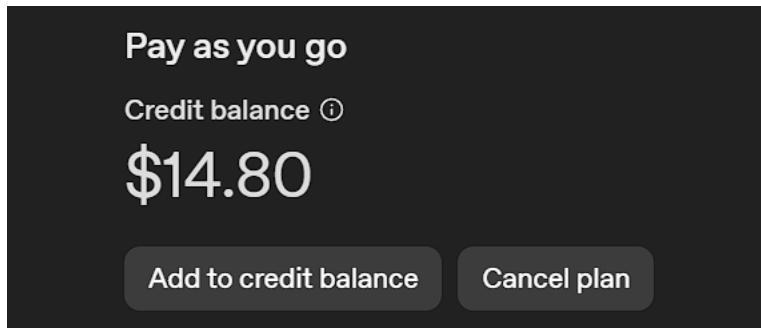
4@| m

All actions taken by the agent—including the triggering indicators, investigative context, and response outcome—are logged for transparency and review. This ensures that every automated decision can be traced, audited, and validated by an analyst.

Step 6: Cost Optimization and Operational Efficiency

Throughout development, cost efficiency was treated as a core design requirement rather than an afterthought. By implementing strict token limits, selecting cost-effective models, and validating usage before execution, the average cost of an investigation was reduced to **well below one cent per run**.





These controls demonstrate how AI-driven security workflows can remain financially viable at scale, even when performing repeated or automated investigations. By combining guardrails with modular model selection, the agent balances analytical depth with predictable, minimal operating costs—making it suitable for continuous SOC usage rather than isolated experimentation.

Summary

By the conclusion of this project, the AI agent is capable of:

- Interpreting natural language security concerns
- Dynamically generating KQL queries
- Analyzing enterprise security telemetry
- Identifying and classifying potential threats
- Applying cost controls and operational guardrails
- Executing automated containment actions through MDE
- Logging findings and responses for analyst review

This project demonstrates how AI agents can be responsibly integrated into SOC workflows—not as a replacement for analysts, but as a **force multiplier** for SOC Teams that accelerates investigations, reduces manual effort, and enables faster, more consistent response to security threats.

From initial detection to automated containment, the agent showcases a practical and scalable approach to modern threat hunting and remediation.